

OpenStreetMap Project Data Wrangling with MongoDB

Ye Xu

Map Area: Boston, MA, United States

<https://mapzen.com/metro-extracts/>

[1. Problems Encountered in the Map](#)

[Over-abbreviated Street Names](#)

[Postal Codes](#)

[2. Data Overview](#)

[3. Additional Ideas](#)

[Contributor statistics and gamification suggestion](#)

[Additional data exploration using MongoDB](#)

[Conclusion](#)

1. Problems Encountered in the Map

I chose the Boston as the city for my project. Initially I looked at a small sample size of the data with the data.py file provisioned in the last lesson of the course. Other than the mappings I finished in the lecture, there are some other abbreviated street types that were not included in the mapping. I iteratively ran this file and added more mappings in the dictionary, finally the mapping dictionary looks something like the follows:

```
mapping = { "St": "Street",
"St.": "Street",
"Ave": "Avenue",
"Rd.": "Road",
"ave": "Avenue",
"avenue": "Avenue",
"rd.": "Road",
"st": "Street",
"street": "Street",
"Street.": "Street",
"Sq.": "Square",
"Pkwy": "Parkway",
"Ct": "Court",
}
```

In the provisioned “data.py” file, I did the following things:

1. For each tag with the value “node” and “way”, the data.py file would process it and convert the tag into a python dictionary.
2. All attributes of "node" and "way" were turned into regular key/value pairs, except the following:
3. Attributes in the CREATED array were added under a key "created"

4. Attributes for latitude and longitude were added to a "pos" array,
5. The values inside "pos" array are floats
6. If second level tag "k" value contains problematic characters, it was ignored
7. If second level tag "k" value starts with "addr:", it was added to a dictionary "address"
8. If second level tag "k" value does not start with "addr:", but contains ":", process it the same as any other tag.
9. If there is a second ":" that separates the type/direction of a street, the tag was ignored.
10. For the address sub field, I used the mapping I created before to store the data.
11. for "way" specifically:
 <nd ref="305896090"/>
 <nd ref="1719825889"/>
 were turned into: "node_refs": ["305896090", "1719825889"]
12. I did each tag element with the above steps, and turn it into a json file element.
13. Then I inserted all of the converted data (.json) into MongoDB.

2. Data Overview

This section contains basic statistics about the dataset and the MongoDB queries used to gather them.

File sizes

```
boston_massachusetts.osm ..... 361 MB  
boston_massachusetts.osm.json.... 412 MB
```

Number of documents

```
>  
db.boston_open.find().count()  
  
1849216
```

Number of nodes

```
> db.boston_open.find({"type": "node"}).count()  
1603318
```

Number of ways

```
> db.boston_open.find({"type": "way"}).count()  
245898
```

Number of unique users

```
> db.boston_open.distinct({"created.user"}).length
336
```

Top 1 contributing user

```
> db.boston_open.aggregate([{"$group":{"_id":"$created.user",
"count":{"$sum":1}}}, {"$sort":{"count":-1}}, {"$limit":1}])
{'ok': 1.0, 'result': [{'_id': 'crschmidt', 'count': 1060051}]}
```

Number of users appearing only once (having 1 post)

```
> db.boston_open.aggregate([{"$group":{"_id":"$created.user",
"count":{"$sum":1}}}, {"$group":{"_id":"$count",
"num_users":{"$sum":1}}}, {"$sort":{"_id":1}}, {"$limit":1}])
{'ok': 1.0, 'result': [{'_id': 1, 'num_users': 215}]}
# “_id” represents postcount
```

3. Additional Ideas

Contributor statistics and gamification suggestion

Another problem in the data is that some nodes already have an “address” attribute. For example, some nodes have “address” = “S Tryon St”. In such a case, there are two options, one is to convert it to the consistent address dictionary, as other nodes, another one is just to leave it like that. I converted it to be consistent with other addresses.

Additional data exploration using MongoDB queries

Top 10 appearing amenities

```
>
db.boston_open.aggregate([{"$match":{"amenity":{"$exists":1}}},
{"$group":{"_id":"$amenity",
"count":{"$sum":1}}}, {"$sort":{"count":1}}, {"$limit":10}])
{'ok': 1.0,
'result': [{'_id': 'parking', 'count': 933},
{'_id': 'bench', 'count': 740},
{'_id': 'school', 'count': 686},
{'_id': 'restaurant', 'count': 446},
{'_id': 'parking_space', 'count': 444},
```

```
{'_id': 'place_of_worship', 'count': 368},
{'_id': 'library', 'count': 324},
{'_id': 'bicycle_parking', 'count': 211},
{'_id': 'cafe', 'count': 170},
{'_id': 'fast_food', 'count': 152}]}
```

Biggest religion (no surprise here)

```
> db.boston_open.aggregate([{"$match":{"amenity":{"$exists":1},
"amenity":"place_of_worship"}},

{"$group":{"_id":"$religion", "count":{"$sum":1}}},

{"$sort":{"count":1}}, {"$limit":1}])

{'ok': 1.0, 'result': [{'_id': 'christian', 'count': 321}]}
```

Most popular cuisines

```
> db.boston_open.aggregate([{"$match":{"amenity":{"$exists":1},
"amenity":"restaurant"}}, {"$group":{"_id":"$cuisine",
"count":{"$sum":1}}}, {"$sort":{"count":1}},
{"$limit":2}])

{'ok': 1.0,
 'result': [{'_id': 'pizza', 'count': 29},
 {'_id': 'american', 'count': 26}]}
```

Conclusion

In this final project I downloaded the Boston area street data map from the open street map website. Before doing any database storing and querying, I cleaned up the raw data in the form described in this report. Then I used mongodb python module to store these data into a mongodb database. These data files are huge, insertion each element into the database took quite a while. But since it is a No-SQL database, the query afterwards was not too slow. Although the data size was as big as 1849214, most of the queries in this report just took less than 5 seconds to finish.